


Alignment of RNA-seq short reads

Vincent Detours
vdetours@ulb.ac.be


1

Mapping RNA-seq reads on the genome




2

Mapping RNA-seq reads on the genome



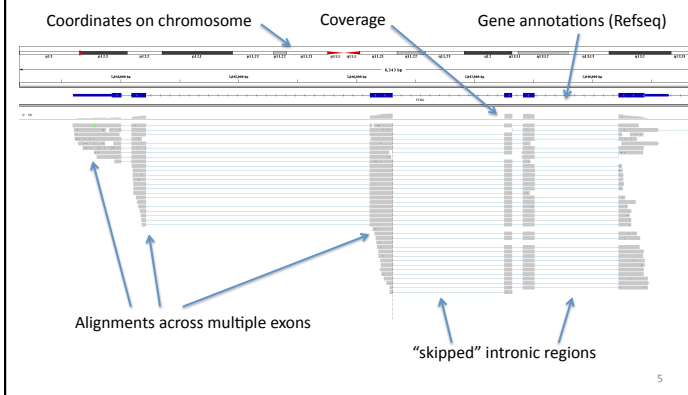
3

Mapping RNA-seq reads on the genome




4

Mapping RNA-seq reads on the genome



Reads are aligned on the human genome...
...and the transcriptome viewed at *any scale*

Gene scale




Counting the reads mapped to a gene gives its expression level in a tractable unit, read per kilobase per million (RPKM), unlike microarrays


6

Reads are aligned on the human genome...
...and the transcriptome viewed at *any scale*

Chromosome scale



Genome scale




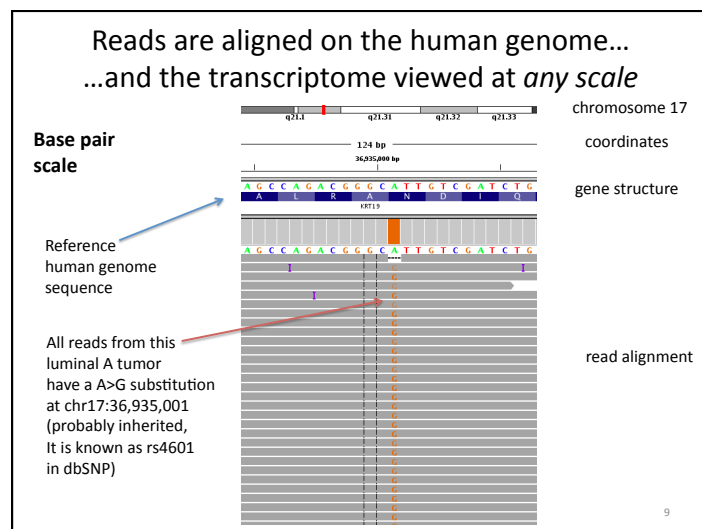
These profiles may be compared to one another or to any positional data set, e.g. evolutionary conservation, transcription factors binding sites, etc.

7

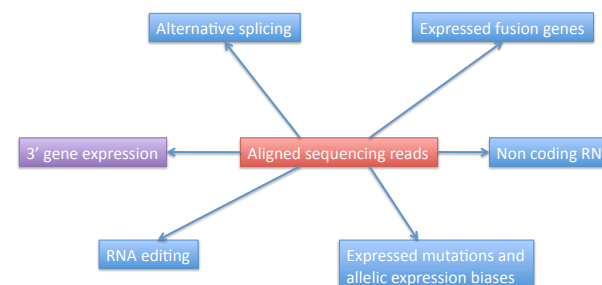
Reads are aligned on the human genome...
...and the transcriptome viewed at *any scale*

Exon scale





Alignment is the prerequisite for a wide range of investigations



Two alignment strategies

De novo assembly:

- The genome structure is not known beforehand
- Short 'shotgun' sequences are assembled together like the piece of a puzzle
- Used for new organism, transcripts, etc.

Alignment on a reference genome

- The genome of the organism has been sequenced, new sequences are placed on the reference sequence
- Use for resequencing (e.g. human cancer, human genetics, etc.)

11

Alignment is challenging

- A typical sequencing run generate millions or billions of reads that need to be aligned on the 3 billions base pair human genome. It's a lot of data and a lot of computations.
- The genome is highly repetitive
 - It evolved as a gigantic copy/paste game through duplications and losses of large and small fragments of DNA. Gene duplication followed by divergence is a major route for genetic innovation.
 - The genome is cluttered with parasitic elements, eg. Alu = 8%, endogenous retroviruses another 8%, etc.
- Sequencing reads are short (2*100bp) and noisy, locating from where they originate in the genome can be error prone, or impossible

12

Aligners speed has dramatically increased thanks to genome indexing

1. Early days : Smith-waterman
2. Genbank rise: BLAST
3. Human genome: BLAT
4. Next generation sequencing: BWA, Bowtie, etc.

2-4 capitalize on the fact that many short queries sequences are aligned on the same long reference sequence.

Dramatic speed-up can be achieved by cleverly indexing the genome... which incurs some loss of accuracy

13

Preprocess the genome: reference indexing

Chr1 : ATGGCTAGCCATTGAAAGGCTATTAC

Chr2 : TGTCAGCGAAGGCTATCTGCTTGAATC

Main idea: creating an index containing the positions of k-mer (here k=5) most often encountered in the reference

GGTCA

chr1:3

chr1:19

chr2:11

TTCGA

chr1:12

chr2:21

Advantage: when mapping a sequence, no need to search the whole the reference to start an alignment (seeding), just check the positions of the corresponding k-mer in the index

14

Seeding: “starting” the alignment

ACCTGCGAAATCGTAAATACTACTTCCTGCGAGCTTGGCCAGCTGCGATGGG


“long” reference sequence (chromosome, mRNA)

AACTGCGATCG

“short” query sequence

15


Seeding: “starting” the alignment



First, use the index to “seed” alignments, that is, find all positions where an alignment can be started.

16

Extending (Smith-Waterman)




Match: +3
Mismatch: -1
Gap opening: -2
Gap extension: -1
parameters of the program

Then, for each “seed”, try to “extend” the alignment as much as possible, to maximize a scoring value

17

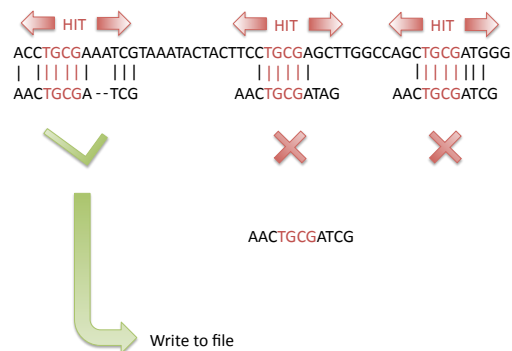
Extending (Smith-Waterman)



Most fast aligners only report the “best” alignment, that is, the one having the highest score


18

Best alignment selection




19

Spliced RNA-seq reads must be aligned on the unspliced genome




20

Incorrect alignments due to pseudogenes




21

Incorrect alignments due to pseudogenes



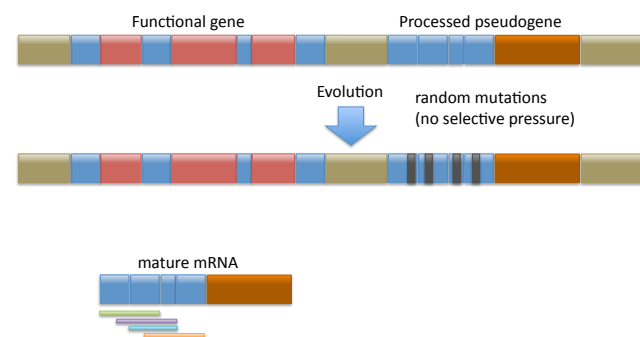
22

Incorrect alignments due to pseudogenes



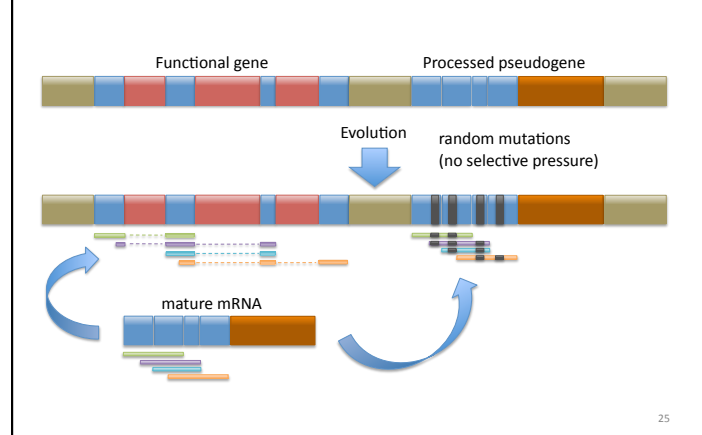
23

Incorrect alignments due to pseudogenes

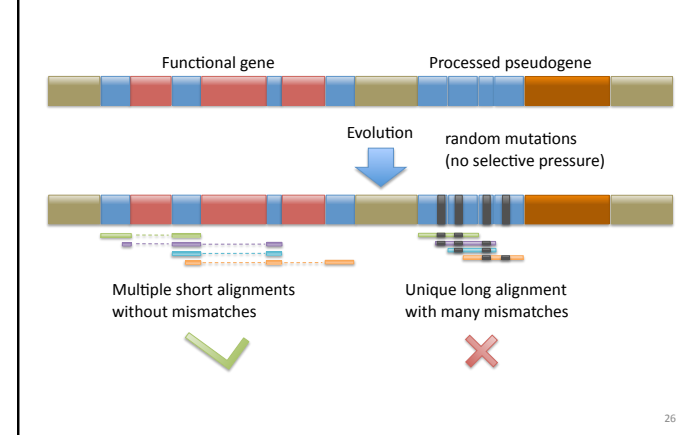


24

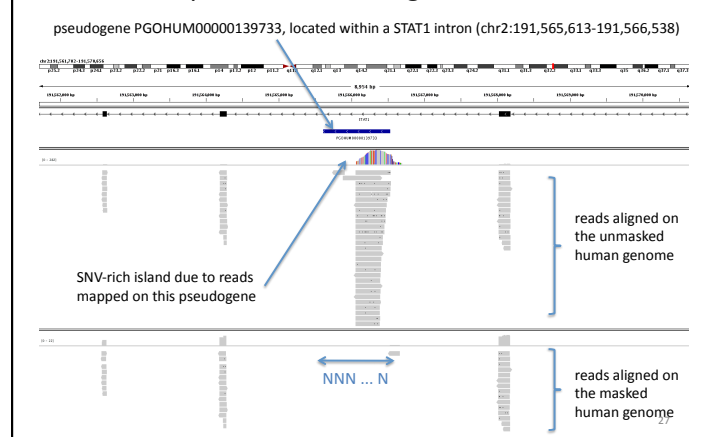
Incorrect alignments due to pseudogenes




Incorrect alignments due to pseudogenes




Example of induced alignment artifact



Common caveats in RNA-seq data




Common caveats in RNA-seq data



Optimal alignment is missed by most short read mappers because the query sequence can be placed as a whole across the intron-exon boundaries (suboptimal alignment due the presence of mismatching bases)

29


Improving alignment accuracy



- A library of splice junction is build on RefSeq, Ensembl and UCSC gene annotations
- This library of splice junctions is added to the reference genome
- Reads are mapped on this custom reference
- Alignment coordinates for reads mapped on junctions are converted to their genomic equivalent

30


Improving alignment accuracy



	ref	start	cigar	type
A1 (genomic)	chr1	p	6M	suboptimal
A2 (junction)	junction chr1 i i+1	p'	6M	optimal (junction)

31

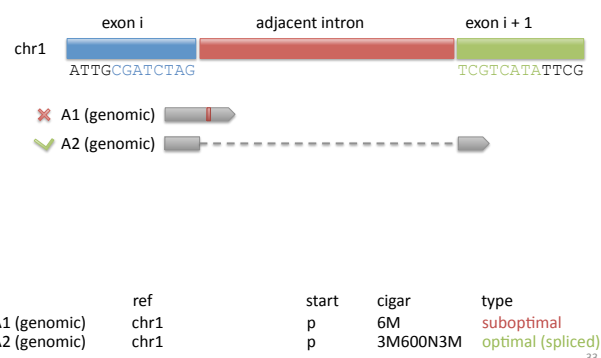
Improving alignment accuracy



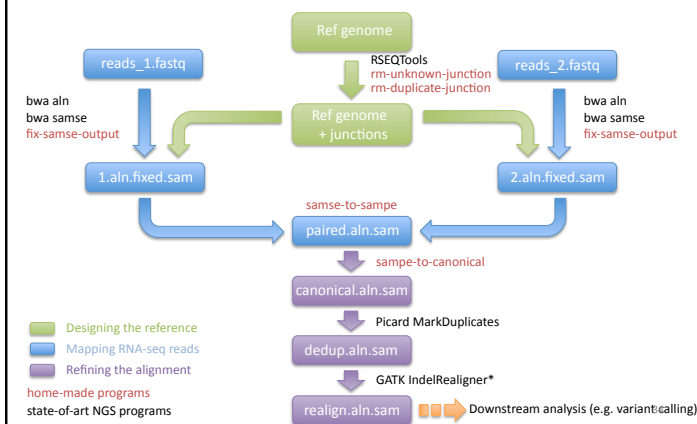
	ref	start	cigar	type
A1 (genomic)	chr1	p	6M	suboptimal
A2 (genomic)	chr1	p	3M600N3M	optimal (spliced)

32

Improving alignment accuracy



Overview of the alignment pipeline



Quantification of gene expression from RNA-seq alignments

- Count the number of reads mapping to each genes
- Divide each gene-wise read count by the length (in kilobases) of gene exons (because longer genes get more reads)
- Divide the gene-wise read counts by the total number millions of reads mapped on the genome (because the deeper you sequence the more read you get)
- The result is expressed in transcripts per millions, a.k.a. TPM
- Many biases remain after this calculation

35

Sources of biases in RNA-seq quantification

- Some highly expressed genes use most of the sequencing depth, e.g. in the thyroid 15% of the reads come from thyroglobulin.
- Hence, other genes seems less expressed
- GC content and other sequence features affect expression measurements...

36